

# Docker

...

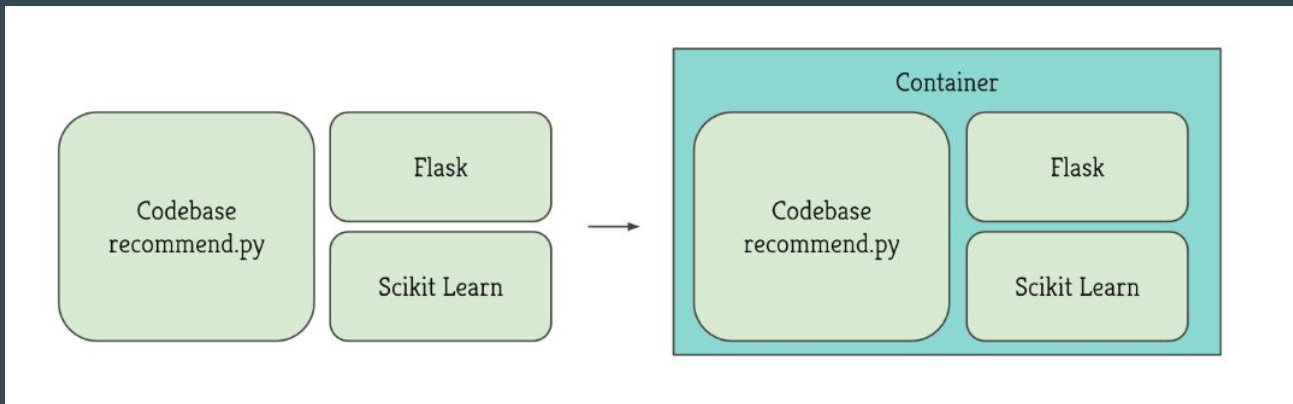
AI Engineering - Recitation 7

# Overview

- Containerization
- Virtual Machines vs Containers
- Docker
- Demo - Docker Images and Containers
- A/B Testing using Containers
- Demo - Load Balancers

# Containerization

- A way to encapsulate or package an executable such that
  - It is isolated from other executables
  - It is easy to move along with its dependencies
  - It is a standalone executable that can be deployed independently
  - It is lightweight in terms of loading and transporting



# Containerization

## Physical Hardware

SE4AI Virtual Machine: 17645-charlie.isri.cmu.edu

Codebase  
recommend-1.py

Flask

Scikit Learn

Ubuntu 14.04

SE4AI Virtual Machine: 17645-titanic.isri.cmu.edu

Codebase  
recommend-2.py

Django

Surprise

Ubuntu 14.04

# Containerization

## Physical Hardware

SE4AI Virtual Machine: 17645-charlie.isri.cmu.edu

### Container

Codebase  
recommend-1.py

Flask

Scikit Learn

Container Runtime Engine

Ubuntu 14.04

SE4AI Virtual Machine: 17645-titanic.isri.cmu.edu

### Container

Codebase  
recommend-2.py

Django

Surprise

Container Runtime Engine

Ubuntu 14.04

# Containerization

Physical Hardware

SE4AI Virtual Machine: 17645-charlie.isri.cmu.edu

Container

Script  
1

Flask

Sklearn

Container

Script  
1

Flask

Sklearn

Container Runtime Engine

Ubuntu 14.04

SE4AI Virtual Machine: 17645-titanic.isri.cmu.edu

Container

Script  
1

Flask

Sklearn

Container

Script  
2

Django

Surprise

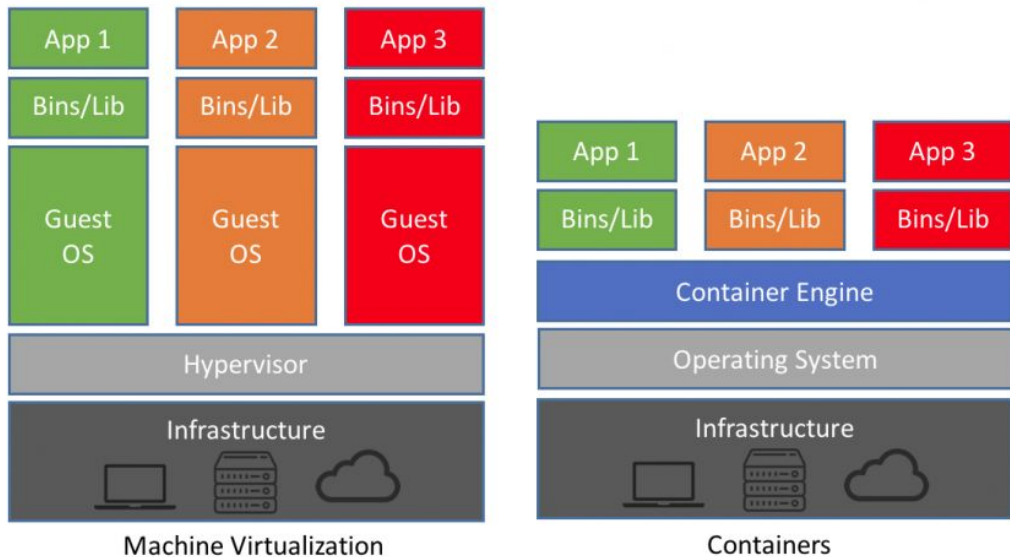
Container Runtime Engine

Ubuntu 14.04

# Virtual Machines vs Containers

- Similarities
  - Both encapsulate your application
- Differences:
  - Size:
    - Containers are smaller in size as they do not contain the OS
  - Portability:
    - VMs are more portable (OS comes along with the VM)
    - Containers are portable as long as the container runtime supports the same format, ie. the same runtime engine has to be installed in the machine where it needs to run
  - OS:
    - Containers are constrained by the OS; VMs are not
    - Each VM has its own OS; Containers share OS of the host machine
  - State:
    - Containers are stateless by default (can be made stateful, although not recommended)

# Virtual Machines vs Containers



Source: <https://blog.netapp.com/blogs/containers-vs-vms/>



# Virtual Machines vs Containers

- Use containers when:
  - You care about the start times of your application (Containers are fast, VMs are slow)
  - Efficiency of resource utilization is of priority (Containers consume less RAM and CPU)
  - You have budget constraints (Docker & Kubernetes are free and open-source)
  - You want to share container images widely (Docker images can be created and shared easily, whereas VM images can be challenging)
- Use VMs when:
  - You are highly concerned with security want to isolate your environment (VMs provide a fully isolated environment by default)
  - You want portability across operating systems (Windows VMs can be deployed on Linux hosts and vice versa; Docker is not as portable)
  - You want to have a rollback feature (VMs can easily go back to a previous snapshot)

# Docker

- Platform-as-a-service product
- Uses OS-level virtualization
- Used to deliver software packages called “containers”
- Terminologies:
  - Image - Everything that is need to configure a fully operational environment
  - Container - A running instance of an image
  - Dockerfile - Definition/Spec to create an image
  - Container Registry - System to host and distribute images
  - Container Repository - Specific physical locations to store related images

# Demo - Docker Images and Containers

- Creating an image using a Dockerfile
- Creating a container using the image
- Running containers on different ports on the same machine
- Inspecting container logs

# A/B Testing using Containers

- A/B testing is an experiment to compare two versions of a variable to find out which performs better in a controlled environment
- Compare performance of different models using this technique
- How to use containers for this?
  - Deploy different containers each having different models
  - Decide a strategy to route users to each model
  - Have a load balancer to execute this strategy
  - Collect results

# Demo - Load Balancers

Outcomes:

- Constructing an efficient load balancer
- Understanding a simple randomizer strategy to route traffic